**Dynamic Pricing System for Urban Parking Lots: Model Analysis and Performance for BHMBCCMKT01**

## 1. Executive Summary

This report details the development and analysis of a sophisticated dynamic pricing system designed for urban parking lots, a capstone project undertaken as part of "Summer Analytics 2025." The core objective of this initiative is to significantly enhance the utilization and revenue generation of urban parking spaces through the implementation of intelligent, data-driven pricing strategies that adapt in real-time to fluctuating market conditions. [1]

Three distinct pricing models were conceptualized, developed, and rigorously evaluated, each progressively increasing in complexity and analytical intelligence. These models — Model 1 (Baseline Linear), Model 2 (Demand-Based Price Function), and Model 3 (Competitive Pricing Model) — illustrate a clear evolution from basic reactive pricing mechanisms to highly sophisticated, market-aware strategic approaches. [1]

A detailed analysis of the pricing behavior specifically for Parking Lot BHMBCCMKT01 has yielded several critical observations regarding model performance. Model 1, designed as a static reference, consistently escalated to a maximum price of $10 and remained largely unresponsive to any subsequent market fluctuations, thereby underscoring the inherent limitations of non-adaptive pricing. In contrast, Model 2, the Demand-Based model, demonstrated a dynamic responsiveness to internal demand factors, maintaining prices that fluctuated smoothly and realistically within a defined range, generally below Model 1's ceiling. Model 3, the Competitive model, exhibited the most advanced real-world applicability by strategically adjusting its prices in direct response to a simulated competitor, showcasing a nuanced and adaptive balance between internal demand pressures and external market forces.

The successful simulation and analysis of this dynamic pricing system affirm its substantial potential to significantly improve urban parking efficiency, enhance customer satisfaction, and optimize revenue management through intelligent, data-driven decision-making.

## 2. Introduction

### 2.1. Background and Motivation

Urban parking spaces represent a critically limited and highly demanded resource within densely populated metropolitan areas. The traditional paradigm of static pricing, where parking rates remain constant irrespective of time of day or prevailing conditions, inevitably leads to significant inefficiencies. Such fixed pricing structures often result in either severe overcrowding during peak demand periods, leading to congestion and lost revenue opportunities, or substantial underutilization during off-peak hours, signifying wasted capacity. [1]

To effectively mitigate these challenges and optimize overall resource utilization, the implementation of dynamic pricing strategies becomes imperative. These strategies involve the continuous adjustment of prices in real-time, informed by fluctuating demand patterns, the competitive landscape of nearby parking alternatives, and prevailing real-time environmental conditions. This project was specifically designed to simulate such a system, with the overarching aim of developing an intelligent, data-driven pricing engine capable of addressing these complex urban challenges. [1]

The necessity of dynamic pricing extends beyond mere technical implementation; it addresses a fundamental economic problem of resource allocation under scarcity. When prices are fixed, they fail to efficiently clear the market, leading to either excess demand (overcrowding) or excess supply (underutilization). By allowing prices to adjust, the system aims to achieve a more optimal equilibrium between the available supply of parking spaces and the fluctuating demand. This mechanism is designed to maximize both the physical utilization of spaces and the potential revenue generated, effectively serving as a market-clearing mechanism. Furthermore, a crucial aspect of this system's design is the requirement for price variations to be "smooth and explainable, not erratic." This stipulation acknowledges that while economic efficiency is a primary driver, market stability and user acceptance are equally vital considerations. Erratic price changes, even if theoretically optimal in a micro-economic sense, could lead to customer dissatisfaction and deter usage, thereby undermining the broader objective of improved utilization and sustained patronage. [1]

### 2.2. Project Objectives

The overarching objective of this capstone project was to design and implement a sophisticated dynamic pricing model for each of 14 simulated urban parking spaces. [1]

The development of these pricing models was guided by several key requirements:

- **Real-time Adaptability:** The developed pricing system was mandated to update prices realistically and continuously in real-time, reflecting the most current operational and environmental conditions. [1]
- **Comprehensive Data Integration:** Price adjustments had to be intelligently informed by a diverse set of real-time data streams. These inputs included historical occupancy patterns, current queue length, nearby traffic congestion levels, indicators for special events (such as holidays), the type of incoming vehicle, and, crucially, competitor parking prices. [1]
- **Base Price Reference:** All pricing models were designed to initiate from a base price of $10. This served as a foundational reference point from which all subsequent dynamic adjustments would be made. [1]
- **Price Stability and Interpretability:** A critical requirement stipulated that price variations must be "smooth and explainable, not erratic." This ensures the system maintains predictability and transparency, fostering user trust and operational stability. [1]
- **Optional Enhancement (Strategic Rerouting):** As an advanced, optional feature, the system was designed to be capable of suggesting rerouting vehicles to nearby alternative parking lots if the current lot became overburdened. This demonstrates a broader, network-level optimization perspective, considering the entire portfolio of parking assets. [1]

### 3. Data Description

#### Data Collection Overview

The foundation of this dynamic pricing system is a comprehensive dataset collected from 14 distinct urban parking spaces. This data spans 73 days, with observations sampled at 18 specific time points per day, occurring at precise 30-minute intervals from 8:00 AM to 4:30 PM. This structured temporal sampling ensures a rich, granular view of daily parking dynamics, providing the necessary resolution for real-time price adjustments. [1]

This 30-minute granularity is crucial for the "real-time" aspect of the dynamic pricing model, enabling rapid adjustments to intra-day fluctuations in demand. It necessitates that the models are designed to process information and update prices efficiently within this timeframe. Furthermore, the reliance on "historical occupancy patterns" implies that these patterns must be intelligently aggregated or summarized to inform predictions for the *next* 30-minute interval, rather than simply reacting to the immediate past. The relatively short daily operational window (8:00 AM to 4:30 PM) suggests that the models are primarily optimized for peak daytime urban parking scenarios, which simplifies the scope by not requiring consideration of overnight or late-night demand patterns. However, it also means the models' direct applicability might be limited to similar operational hours. [1]

#### Detailed Feature Breakdown

Each recorded instance within the dataset provides a rich set of features, reflecting the state of each parking lot at a given time step. These features serve as critical inputs for the dynamic pricing models: [1]

- **Location Information:**
  - **Latitude and Longitude:** These geographic coordinates are vital for calculating the precise proximity of a given parking space to its competitors. This spatial intelligence forms a cornerstone for the Competitive Pricing Model (Model 3), enabling the system to understand its market position relative to other parking options.
- **Parking Lot Features:**
  - **Capacity:** Represents the maximum number of vehicles that can be accommodated in a parking lot. This metric is fundamental for calculating the occupancy rate, a key indicator of utilization.
  - **Occupancy:** The current number of vehicles parked. This is a primary, real-time indicator of current demand and the immediate utilization level of the lot.

- **Queue Length:** The number of vehicles currently waiting for entry into the parking lot. This serves as a direct and immediate measure of unmet demand and potential congestion, signaling high pressure on the lot and justifying potential price adjustments.

- **Vehicle Information:**
  - **Type of Incoming Vehicle:** Categorizes incoming vehicles as car, bike, or truck. This feature allows for the potential differentiation of demand or pricing based on vehicle characteristics, such as the space required or typical duration of stay, which can influence a lot's overall capacity utilization.

- **Environmental Conditions:**
  - **Nearby Traffic Congestion Level:** An external environmental factor that can significantly influence the accessibility of the parking lot and the perceived urgency or value of available spaces. Higher congestion might increase the perceived value of an available spot.
  - **Special Day Indicator:** A binary flag (e.g., holiday, event day) that captures exogenous demand shocks. Special days typically lead to significant deviations from normal demand patterns, requiring adaptive pricing to capitalize on or manage these surges.

The interplay of these diverse features ensures that demand for parking spaces fluctuates realistically throughout the day, underscoring the necessity and complexity of a robust dynamic pricing approach. [1]

## 4. Pricing Model Design and Implementation

This section elucidates the architectural design and underlying logic of the three developed pricing models, highlighting their progressive sophistication and how each addresses specific project objectives.

### 4.1. Model 1: Baseline Linear Model

Model 1 is designed as a foundational, simple reactive system. Its pricing mechanism is based on a linear relationship where the next price ($Price_{t+1}$) is a direct function of the previous price ($Price_t$) and the current occupancy rate. The formula provided is: $Price_{t+1} = Price_t + \alpha * (Occupancy / Capacity)$. [1] This model acts as a straightforward reference point, illustrating the most basic form of price adjustment based on utilization.

While its simplicity offers ease of understanding and implementation, this approach also represents a significant limitation. Its sole reliance on occupancy and a linear adjustment means it fails to account for other critical demand factors such as queue length, external traffic conditions, special events, or the crucial aspect of competitor pricing. This inherent inflexibility leads to its observed static behavior once a certain occupancy threshold is met, effectively hitting a price ceiling (likely the $10 base price or a predefined maximum). This makes it less "dynamic" in a comprehensive sense. The rapid rise to $10, as observed in the accompanying graph for Parking Lot BHMBCCMKT01, suggests that for this specific lot, occupancy quickly reached a level that, according to Model 1's alpha coefficient and initial price, triggered the maximum permissible price. This behavior underscores why it serves as a "reference point"; it starkly contrasts with the more adaptive models by demonstrating the inefficiencies of a system that cannot dynamically adjust downwards when demand wanes, potentially leading to underutilization if demand is low but prices remain high. [1]

Model 1's primary utility lies in its function as a benchmark. By observing its straightforward, occupancy-driven behavior, it provides a clear contrast against which the dynamism, sophistication, and real-world applicability of Model 2 and Model 3 can be quantitatively and qualitatively evaluated. It effectively highlights the inefficiencies inherent in non-adaptive pricing strategies.

### 4.2. Model 2: Demand-Based Price Function

Model 2 represents a significant leap in sophistication, moving beyond simple occupancy to incorporate a comprehensive mathematical demand function. This function is designed to provide a more nuanced and accurate understanding of real-time demand by integrating multiple influencing factors. [1]

The model constructs a mathematical function that intelligently aggregates the influence of various demand drivers. This involves assigning weights or coefficients to each factor to reflect its hypothesized impact on overall demand. The key features incorporated into this demand function include: Occupancy rate, Queue length, Traffic level, Special day indicator, and Vehicle type. [1] The hypothesized influence of each feature is critical: for instance, higher occupancy and longer queues generally indicate higher immediate demand, while a "Special Day" flag signifies an exogenous surge in demand. Traffic levels might have a more complex impact, potentially increasing demand for convenient parking if congestion is high, or decreasing it if access is severely impeded.

The problem statement encouraged the design and implementation of "more effective and sophisticated pricing strategies" beyond the "linear and simple baselines" provided as examples. This implies that the actual implementation for Model 2 likely goes beyond a simple linear combination of features. A more sophisticated approach might involve non-linear transformations of input features (e.g., logarithmic scaling for queue length to reflect diminishing returns of impact), interaction terms between features (e.g., the impact of traffic might be different on a special day), or even a more complex, albeit still numpy/pandas-based, statistical or machine learning approach to derive the demand value. For instance, VehicleTypeWeight could be implemented as a lookup table or a set of one-hot encoded features with specific coefficients, reflecting differential demand or space requirements. The observed "smooth and bounded" fluctuations in the graph for Model 2 are a testament to the successful tuning of the Lambda parameter and the demand normalization process, indicating that even with potentially sophisticated underlying logic, the model effectively manages to prevent erratic price changes, successfully balancing responsiveness with stability, a key project objective. [1]

The calculated demand value is normalized to ensure consistency across different scales of input features and to prevent disproportionate influence from any single factor. This normalization is crucial for maintaining price stability. The final price adjustment is then applied using the formula: Price = Base Price * (1 + $\lambda$ * NormalizedDemand). [1] A critical requirement for Model 2 is that price variations must be "smooth and bounded," specifically "not more than 2x or less than 0.5x base."

Given the $10 base price, this translates to an expected price range of $5 to $20. The graph demonstrates that Model 2 operates effectively within this range, predominantly fluctuating between approximately $5.20 and $6.70 for the observed parking lot. [1]

The features influencing the demand function and their hypothesized impact are summarized in Table 1, providing a clear explanation of the model's core inputs and underlying assumptions.

**Table 1: Demand Function Features and Hypothesized Impact**

| Feature | Description | Hypothesized Impact on Demand (and thus Price) |
|---|---|---|
| Occupancy Rate | Ratio of occupied spaces to total capacity | Higher occupancy generally indicates higher demand, leading to an increase in price. |
| Queue Length | Number of vehicles currently waiting for entry | A longer queue signifies immediate, unmet demand, justifying a higher price. |
| Traffic Level | Level of nearby road congestion | Higher traffic can increase the perceived value of an available parking spot, thus increasing demand and price, especially for convenient locations. |
| Special Day | Binary indicator for holidays or special events | Special days typically lead to a surge in demand, warranting an increase in price. |
| Vehicle Type | Type of incoming vehicle (e.g., car, bike, or truck) | Different vehicle types may have varying space requirements or typical parking durations, influencing demand and potentially leading to differentiated pricing weights. |

**4.3. Model 3: Competitive Pricing Model**

Model 3 represents the pinnacle of sophistication in this project, integrating "location intelligence" and simulating "real-world competition" into the dynamic pricing strategy. This model moves beyond solely considering internal demand factors to actively account for the external market landscape. [1]

A foundational step for Model 3 involves leveraging the provided latitude and longitude data for each parking space. These geographic coordinates are used to accurately calculate the proximity of nearby parking spaces, which are then identified as potential competitors. [1]

The core intelligence of Model 3 lies in its ability to factor competitor prices into its own pricing decisions. It employs sophisticated competitive logic based on two primary scenarios: [1]

- **Scenario 1: Current lot is full and nearby competitor lots are cheaper.** In this situation, the model can either suggest rerouting vehicles to the cheaper, less-utilized nearby lots (a strategic network-level decision) or, if appropriate, reduce the current lot's price. The latter might seem counterintuitive for a full lot but could be a strategy to clear a queue more quickly or to maintain attractiveness for future demand if the 'full' state is temporary.
- **Scenario 2: Nearby competitor lots are expensive.** When competitors are pricing higher, Model 3 identifies an opportunity to increase its own price while still remaining attractive and competitive to potential customers. This allows the lot to capture additional revenue without losing market share.

This model is not merely reactive to internal demand; it engages in a strategic market game. Its objective implicitly balances revenue maximization with maintaining competitive attractiveness and optimizing utilization within a multi-lot ecosystem. The observed strong correlation with competitor prices in the graph directly validates this strategic behavior, demonstrating the model's ability to adapt to external market dynamics. The "rerouting" suggestion, even if not fully implemented in the pricing mechanism itself, represents a crucial piece of "business thinking." It implies a potential for network-level optimization, where the system considers the entire portfolio of parking lots rather than just isolated units. This consideration can influence pricing decisions by reducing the immediate pressure to drastically lower prices when full, as alternative solutions (rerouting) might be available. [1]

This model explicitly encourages "creativity and business thinking along with technical skills," highlighting that its development requires not just coding prowess but also a deep understanding of market dynamics, competitive strategy, and customer behavior. [1]

A comparative analysis of the three pricing models, highlighting their primary inputs, core logic, responsiveness, complexity, realism, and objectives, is provided in Table 2.

**Table 2: Comparative Analysis of Pricing Models**

| Feature | Model 1 (Baseline Linear) | Model 2 (Demand-Based) | Model 3 (Competitive) |
|---|---|---|---|
| **Primary Inputs** | Previous Price, Occupancy Rate | Occupancy Rate, Queue Length, Traffic Level, Special Day, Vehicle Type | All Model 2 inputs + Competitor Prices, Geographic Proximity |
| **Core Logic** | Simple linear price increase based on occupancy; serves as a static reference point. | Price adjusted dynamically based on a comprehensive, multi-factor mathematical demand function. | Strategic pricing that considers both internal demand factors and external competitor prices, aiming for market advantage. |
| **Responsiveness** | Limited; quickly becomes static once a certain | Highly responsive and adaptive to real-time | Highly responsive to both internal demand and the external |

| | | | |
|---|---|---|---|
| | occupancy threshold is met. | fluctuations in internal demand drivers. | competitive landscape, demonstrating strategic market awareness. |
| **Complexity** | Low (simple, rule-based) | Moderate (multi-variable function, normalization) | High (incorporates spatial intelligence, strategic decision-making) |
| **Realism** | Low (does not reflect market dynamics) | Moderate (captures demand-side dynamics) | High (simulates real-world competitive market behavior) |
| **Objective** | To provide a fundamental benchmark for comparison against more advanced models. | To optimize parking lot utilization and revenue by accurately reflecting real-time demand. | To optimize utilization and revenue by strategically positioning prices within a competitive market. |

## 5. Real-Time Simulation and Visualization

### 5.1. Real-Time Simulation with Pathway

The project leverages Pathway, a real-time data processing framework, as an integral component for simulating the dynamic pricing system. Pathway facilitates the continuous ingestion and processing of data streams, enabling the models to operate in a live environment. [1]

Key capabilities utilized from Pathway include:

- **Data Streaming and Order Preservation:** Pathway is utilized for ingesting data streams with simulated delays, critically preserving the original timestamp order. This ensures that the pricing models receive information in a chronologically accurate sequence, vital for real-time decision-making. [1]
- **Real-time Feature Processing:** The framework allows for the continuous processing of various input features (e.g., occupancy, queue length, traffic levels) as they arrive, ensuring that the pricing models always operate on the most current data. [1]
- **Continuous Prediction Emission:** Pathway enables the continuous emission of pricing predictions from the models, allowing for dynamic adjustments to be pushed out in real-time, fulfilling a core requirement of the project. [1]

All code for the project, including the Pathway integration and pricing logic, was developed and executed within the Google Colab environment. This cloud-based platform provided the necessary computational resources and collaborative framework. [1]

### 5.2. Visualization Requirements with Bokeh

Bokeh, a powerful interactive visualization library, was selected to fulfill the project's visualization requirements. It enables the creation of real-time, dynamic plots that are crucial for monitoring and understanding the pricing behavior of the developed models. [1]

A key mandate for the visualizations is to "visually justify pricing behavior." This means the plots are not merely descriptive; they serve as direct evidence to explain *why* prices change as they do, linking observed patterns back to the underlying model logic and real-time data inputs. This visual justification is paramount for demonstrating the models' adherence to project objectives like smoothness and responsiveness. The explicit requirement for Bokeh visualizations to visually justify pricing behavior, coupled with the provision of the "Dynamic Pricing for Parking Lot: BHMBCCMKT01" graph, indicates that the visual output is a primary means of evaluation. The graph is the central piece of evidence for assessing the models against qualitative objectives, such as "smooth and explainable, not erratic" price variation. It allows for a direct, qualitative assessment of how responsive, stable, and strategically aligned each model is. The visual comparison of multiple models and a competitor price on a single plot is highly effective for rapid comparative analysis. Given that the project "has no objective metric for evaluation like in a hackathon," the visual evidence becomes the cornerstone of the project's assessment. The ability to clearly demonstrate model behavior and its rationale through graphs is thus as important as the underlying code. [1]

The project specifically suggested and produced real-time pricing line plots for each parking space, alongside comparative plots that illustrate the developed models' prices against simulated competitor prices. The provided graph for Parking Lot BHMBCCMKT01 is a direct example of such a critical visualization. [1]

## 6. Results and Analysis

This section provides a detailed interpretation of the provided "Dynamic Pricing for Parking Lot: BHMBCCMKT01" graph, meticulously linking the observed pricing behaviors of each model to their underlying design principles and the overarching project objectives.

### Graph Overview

The graph, titled "Dynamic Pricing for Parking Lot: BHMBCCMKT01," illustrates the price trajectories of the three developed models—Model 1 (Baseline), Model 2 (Demand-Responsive), and Model 3 (Competitive)—over approximately 320 time steps, alongside a simulated competitor price. This visualization is crucial for understanding the dynamic interactions and performance of each pricing strategy for a specific parking lot.

The graph's title explicitly states "Dynamic Pricing for Parking Lot: BHMBCCMKT01." This means the observed pricing behaviors are unique to *this* particular parking lot and may not be universally representative of all 14 lots in the dataset. The analysis must acknowledge that the specific demand patterns, competitive dynamics, and resulting price fluctuations seen in this graph are localized to BHMBCCMKT01. For example, the fact that Model 2 and 3 consistently price below $7 (and below Model 1's $10 ceiling) suggests that this specific lot might experience generally lower average demand, higher local competition, or a lower optimal price point compared to other lots. If the demand for BHMBCCMKT01 is indeed often lower, Model 1's behavior (rapidly hitting $10) further highlights its insensitivity to actual market conditions. It implies that Model 1 might be reacting to a simple occupancy threshold that is easily met, rather than true demand, making it inefficient for lots with varying demand profiles.

**6.1. Analysis of Model 1 (Baseline Linear)**

Model 1 initiates at approximately $5.20 and exhibits a rapid, linear increase in price. By around time step 120, its price stabilizes precisely at $10.00, remaining flat for the remainder of the simulation. This behavior is entirely consistent with its design as a simple linear model: Price_t+1 = Price_t + α * (Occupancy / Capacity). [1] The swift ascent suggests that the parking lot's occupancy quickly reached a threshold that, combined with the model's

α coefficient, drove the price upwards until it hit a predefined ceiling or the explicit $10 base price, which likely acts as an upper bound in this model. The most significant limitation of Model 1 is its static nature after reaching the $10 ceiling. It completely loses its dynamism and responsiveness to any subsequent changes in occupancy, queue length, traffic, or competitor prices. While serving its purpose as a "reference point" [1], it visually demonstrates the inherent inefficiencies of a non-adaptive pricing strategy in a dynamic urban environment.

**6.2. Analysis of Model 2 (Demand-Responsive)**

Model 2, starting around $5.20, displays significantly more dynamic fluctuations than Model 1. Its price generally oscillates between approximately $5.20 and $6.70. While dynamic, its movements appear relatively smooth, without erratic spikes or drops. This dynamic yet smooth behavior is a direct reflection of its sophisticated mathematical demand function. This model continuously calculates demand based on multiple real-time factors (occupancy, queue length, traffic, special day, vehicle type) and adjusts the price accordingly using the formula Price = Base Price * (1 + λ * NormalizedDemand). [1] The observed rises correspond to periods of higher calculated demand, while dips indicate lower demand. Its generally lower price range compared to Model 1's $10 ceiling suggests that for Parking Lot BHMBCCMKT01, the demand-driven optimal price is often below the static baseline, indicating that the market for this specific lot might not support consistently high prices. Model 2 successfully adheres to the crucial project objective of ensuring "smooth and explainable, not erratic" price variations. [1] Its controlled fluctuations demonstrate that the demand normalization and price adjustment mechanisms are well-tuned to provide responsiveness without sacrificing stability or predictability.

### 6.3. Analysis of Model 3 (Competitive)

Model 3, starting around $6.00, exhibits the most dynamic and strategically intelligent behavior. Its price fluctuations (ranging roughly between $5.50 and $6.70) show a clear and strong correlation with the "Simulated Competitor Price" (dashed red line). Model 3 often mirrors the competitor's movements, frequently pricing itself slightly below or at a similar level. This strong correlation is compelling evidence that Model 3 effectively incorporates competitive intelligence into its pricing decisions. As per its design, it actively calculates geographic proximity to competitors and factors their prices into its own strategy. [1]

When the simulated competitor price rises (e.g., around time step 120-130 or 270-280), Model 3 tends to follow suit, capitalizing on the higher market tolerance. Conversely, when the competitor's price drops (e.g., around time step 170-180 or 230-240), Model 3 also adjusts downwards, demonstrating a strategy to maintain attractiveness and prevent customer loss. Model 3 exemplifies advanced business thinking. It aims not only to respond to internal demand but also to strategically position itself within the market. Its ability to adapt to external competitive pressures makes it the most realistic and robust model for optimizing both revenue and utilization in a multi-lot, competitive urban environment.

### 6.4. Comparative Performance and Overall Insights

The comparison clearly highlights Model 1's static and limited nature against the true dynamism of Model 2 and Model 3. Model 2 and Model 3 are significantly more aligned with the project's core objective of developing an intelligent, real-time, and adaptive pricing engine. It is notable that for Parking Lot BHMBCCMKT01, Model 2 and Model 3 consistently operate at a lower price point than Model 1's ceiling of $10. This suggests that the optimal demand-driven or competitively strategic price for this specific lot is often below the initial $10 base price, indicating a market where lower prices might be necessary to attract or retain customers. The $10 base price serves effectively as an anchor or starting point, allowing for significant downward adjustments based on market conditions. Despite their varying levels of dynamism, all three models exhibit relatively smooth price transitions. This confirms successful implementation of the "smooth and explainable, not erratic" requirement, ensuring that the pricing behavior is logical and predictable, which is crucial for both operational stability and user acceptance.

## 7. Discussion and Conclusion

### 7.1. Summary of Key Findings

This project successfully demonstrates the development and implementation of three progressive dynamic pricing models for urban parking lots. Model 1 served as an essential baseline, clearly illustrating the limitations of static pricing. Model 2 effectively showcased the power of demand-responsive pricing by integrating multiple real-time features to derive nuanced price adjustments. Model 3 further advanced the system by incorporating competitive intelligence, strategically adapting prices in response to market dynamics. The visual analysis of Parking Lot BHMBCCMKT01 confirmed that Model 2 and Model 3 are significantly more effective in achieving the project's objectives of optimizing parking lot utilization and revenue through adaptive, data-driven pricing, while maintaining price stability.

### 7.2. Assumptions Made

The effectiveness of the developed dynamic pricing system relies on several underlying assumptions:

- **Data Accuracy and Timeliness:** A fundamental assumption is that the "real-time data streams" [1] provided are consistently accurate, complete, and delivered with sufficient timeliness (e.g., within the 30-minute interval) to enable effective and responsive price adjustments. Any significant delays or inaccuracies in data ingestion could compromise the models' performance and lead to suboptimal pricing decisions.
- **Model Parameters:** The effectiveness of Model 2 and Model 3 relies on the appropriate tuning of various parameters, such as the coefficients (α, β, γ, δ, ε) within the demand function and the price elasticity factor (λ). It is assumed that these parameters have been calibrated to reflect realistic market responses and optimize the desired pricing behavior for the specific parking lot.
- **Competitor Behavior (for Model 3):** For Model 3, it is assumed that the "Simulated Competitor Price" accurately reflects realistic competitor pricing strategies and that competitor price data is consistently observable and reliable. The model's strategic decisions are contingent on the validity and representativeness of this competitive intelligence.
- **Market Rationality:** The underlying economic theory implicitly assumes that parking users respond rationally to price signals and demand cues (e.g., lower prices attract more users, higher prices deter them, and queues indicate high demand). Deviations from this rational behavior, such as brand loyalty or lack of information, could impact the models' effectiveness in real-world scenarios.

### 7.3. Strengths and Limitations

The project exhibits several notable strengths and inherent limitations:

**Strengths:**

- **Comprehensive Feature Utilization:** The models effectively leverage a wide array of real-time data features (occupancy, queue, traffic, special events, vehicle type, competitor prices) to inform pricing decisions, leading to sophisticated and adaptive strategies that reflect complex market dynamics.

- **Modular and Progressive Design:** The structured progression from a simple baseline to complex demand-responsive and competitive models demonstrates a robust and scalable development approach, allowing for incremental sophistication.
- **Real-Time Capability:** Successful integration with Pathway for real-time data ingestion, processing, and continuous price emission validates the system's ability to operate in a dynamic environment, crucial for urban parking management.
- **Visual Justification:** The effective use of Bokeh for real-time visualizations provides clear, intuitive justification for the pricing behavior, making the models' logic transparent and verifiable, which is particularly important given the absence of a single objective metric.

**Limitations:**

- **Generalizability:** The detailed analysis and observed behaviors are specific to Parking Lot BHMBCCMKT01. The optimal pricing strategies and model performance might vary significantly across the other 13 parking spaces due to differing local demand profiles, competitive landscapes, or operational characteristics. Further analysis across all lots would be required for broader conclusions.
- **Absence of Objective Metric:** As noted in the problem statement, the project lacked an explicit objective metric for evaluation (e.g., maximizing total revenue or utilization over a period). [1] This means performance assessment is primarily qualitative, relying on visual analysis and adherence to design principles rather than quantifiable optimization of a specific business outcome.
- **"From Scratch" Constraint:** The core requirement to implement all pricing models "from scratch" using "only numpy, pandas libraries" [1] restricted the use of more advanced machine learning frameworks (e.g., scikit-learn, TensorFlow, PyTorch). While fostering fundamental understanding, this constraint might limit the ultimate sophistication or predictive power of the demand models compared to what could be achieved with more advanced tools.
- **Simplified Competitive Model:** While Model 3 incorporates competition, it assumes a "Simulated Competitor Price." Real-world competitive dynamics can be far more complex, involving strategic interactions, price wars, and non-price competition that are not fully captured by a single simulated price feed.

### 7.4. Future Enhancements and Extensions

Building upon the foundation established by this project, several avenues for future enhancement and extension are identified:

- **Explicit Revenue/Utilization Optimization:** Integrate an explicit objective function (e.g., total revenue, average utilization, or profit margin) to quantitatively evaluate and further tune the models. This would allow for data-driven optimization of pricing parameters towards specific business goals.

- **Predictive Demand Modeling:** Enhance the demand function by incorporating time-series forecasting techniques or more advanced predictive machine learning models. This would allow the system to anticipate future demand patterns, rather than solely reacting to current conditions, enabling proactive pricing strategies.
- **Reinforcement Learning for Pricing:** Explore the application of reinforcement learning (RL) frameworks. An RL agent could learn optimal pricing strategies through continuous interaction with the simulated environment, adapting its policies based on observed outcomes (e.g., revenue generated, spaces utilized) over time, potentially discovering non-obvious optimal strategies.
- **Network-Level Optimization:** Fully implement and optimize the "rerouting" suggestion across all 14 parking lots. This would transform the system into a network-wide optimization strategy, balancing demand and supply across multiple locations for collective efficiency and potentially higher overall system revenue.
- **Nuanced User Behavior Modeling:** Incorporate more sophisticated models of user behavior, including price elasticity of demand for different user segments, varying willingness-to-pay, and the psychological impact of factors like queue length or traffic congestion on decision-making.
- **Adaptive Parameter Tuning:** Develop mechanisms for dynamically adjusting model parameters (e.g., $\alpha$, $\lambda$ coefficients, or weights in the demand function) in real-time based on observed market responses and performance metrics, moving beyond static parameter settings. This would allow the models to adapt to changing market conditions over longer periods.
- **A/B Testing Framework:** For real-world deployment, design an A/B testing framework to compare the performance of different pricing strategies in live scenarios. This would allow for continuous improvement and validation of the models' effectiveness in real-world conditions.